

Assignment 8: Time Series Analysis

Rachel Williams

Fall 2024

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on generalized linear models.

Directions

1. Rename this file `<FirstLast>_A08_TimeSeries.Rmd` (replacing `<FirstLast>` with your first and last name).
2. Change “Student Name” on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure to **answer the questions** in this assignment document.
5. When you have completed the assignment, **Knit** the text and code into a single PDF file.

Set up

1. Set up your session:
 - Check your working directory
 - Load the tidyverse, lubridate, zoo, and trend packages
 - Set your ggplot theme

```
library(here)
```

```
## here() starts at /home/guest/EDE_Fall2024
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(lubridate)
library(zoo)
```

```
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
```

```
library(trend)
library(dplyr)
here()
```

```
## [1] "/home/guest/EDE_Fall2024"
```

2. Import the ten datasets from the Ozone_TimeSeries folder in the Raw data folder. These contain ozone concentrations at Garinger High School in North Carolina from 2010-2019 (the EPA air database only allows downloads for one year at a time). Import these either individually or in bulk and then combine them into a single dataframe named **GaringerOzone** of 3589 observation and 20 variables.

```
#1
Ozone1 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2010_raw.csv"))
Ozone2 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2011_raw.csv"))
Ozone3 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2012_raw.csv"))
Ozone4 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2013_raw.csv"))
Ozone5 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2014_raw.csv"))
Ozone6 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2015_raw.csv"))
Ozone7 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2016_raw.csv"))
Ozone8 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2017_raw.csv"))
Ozone9 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2018_raw.csv"))
Ozone10 <- read.csv(here("Data/Raw/Ozone_TimeSeries/EPAair_03_GaringerNC2019_raw.csv"))
GaringerOzone <- rbind(Ozone1,Ozone2,Ozone3,Ozone4,Ozone5,Ozone6,Ozone7,Ozone8,
                      Ozone9,Ozone10)
```

Wrangle

3. Set your date column as a date class.
4. Wrangle your dataset so that it only contains the columns Date, Daily.Max.8.hour.Ozone.Concentration, and DAILY_AQI_VALUE.
5. Notice there are a few days in each year that are missing ozone concentrations. We want to generate a daily dataset, so we will need to fill in any missing days with NA. Create a new data frame that contains a sequence of dates from 2010-01-01 to 2019-12-31 (hint: `as.data.frame(seq())`). Call this new data frame Days. Rename the column name in Days to "Date".
6. Use a `left_join` to combine the data frames. Specify the correct order of data frames within this function so that the final dimensions are 3652 rows and 3 columns. Call your combined data frame **GaringerOzone**.

```

# 3
GaringerOzone$Date <- as.Date(GaringerOzone$Date, format="%m/%d/%Y")
# 4
GaringerOzoneProcessed <- GaringerOzone %>%
  select(Date, Daily.Max.8.hour.Ozone.Concentration, DAILY_AQI_VALUE)

# 5
Days <- as.data.frame(seq(from=as.Date("2010-01-01"), to=as.Date("2019-12-31"), by="day"))
colnames(Days) <- "Date"
# 6
GaringerOzone <- left_join(Days, GaringerOzoneProcessed, by="Date")

```

Visualize


7. Create a line plot depicting ozone concentrations over time. In this case, we will plot actual concentrations in ppm, not AQI values. Format your axes accordingly. Add a smoothed line showing any linear trend of your data. Does your plot suggest a trend in ozone concentration over time?

```

#7
ggplot(GaringerOzoneProcessed, aes(x=Date, y=Daily.Max.8.hour.Ozone.Concentration))+
  geom_line() +geom_smooth(method=lm)

```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



Answer: My plot suggest a slightly negative trend in ozone concentration over time during the 2010s.

Time Series Analysis

Study question: Have ozone concentrations changed over the 2010s at this station?

8. Use a linear interpolation to fill in missing daily data for ozone concentration. Why didn't we use a piecewise constant or spline interpolation?

#8

```
Garinger_data_clean <- GaringerOzone %>%  
  mutate( Clean.Ozone = zoo::na.approx(Daily.Max.8.hour.Ozone.Concentration))
```

Answer: We didn't use a piecewise constant because it copies the closest value instead of continuing the linear pattern. If there were a lot of N/As in a row it might change the slope of the linear trend. We didn't use spline interpolation because we are looking for linear trends and not quadratic.

9. Create a new data frame called `GaringerOzone.monthly` that contains aggregated data: mean ozone concentrations for each month. In your pipe, you will need to first add columns for year and month to form the groupings. In a separate line of code, create a new Date column with each month-year combination being set as the first day of the month (this is for graphing purposes only)

#9

```
GaringerOzone.monthly <- Garinger_data_clean %>%  
  mutate(month= month(Date)) %>%  
  mutate(year= year(Date)) %>% group_by(year, month) %>%  
  summarize(mean_concentration = mean(Clean.Ozone, na.rm = TRUE), .groups = 'drop') %>%  
  mutate(Date = as.Date(paste(year, month, "01", sep = "-")))  
  
class(GaringerOzone.monthly$Date)
```

```
## [1] "Date"
```

10. Generate two time series objects. Name the first `GaringerOzone.daily.ts` and base it on the dataframe of daily observations. Name the second `GaringerOzone.monthly.ts` and base it on the monthly average ozone values. Be sure that each specifies the correct start and end dates and the frequency of the time series.

#10


```
f_month <- month(first(Garinger_data_clean$Date))  
f_year <- year(first(Garinger_data_clean$Date))  
  
f_month2 <- month(first(GaringerOzone.monthly$Date))  
f_year2 <- year(first(GaringerOzone.monthly$Date))  
  
GaringerOzone.daily.ts <- ts(Garinger_data_clean$Clean.Ozone,  
  start=c(f_year,f_month), frequency=365)
```

```
GaringerOzone.monthly.ts <- ts(GaringerOzone.monthly$mean_concentration,
                              start=c(f_year2, f_month2), frequency=12)
```


11. Decompose the daily and the monthly time series objects and plot the components using the `plot()` function.

```
#11
Dailydecompose <- stl(GaringerOzone.daily.ts, s.window="periodic")
Monthlydecompose <- stl(GaringerOzone.monthly.ts, s.window="periodic")

plot(Dailydecompose)
```



```
plot(Monthlydecompose)
```



12. Run a monotonic trend analysis for the monthly Ozone series. In this case the seasonal Mann-Kendall is most appropriate; why is this?

#12

```
monthly_trend <- Kendall::SeasonalMannKendall(GaringerOzone.monthly.ts)
summary(monthly_trend)
```

```
## Score = -77 , Var(Score) = 1499
## denominator = 539.4972
## tau = -0.143, 2-sided pvalue =0.046724
```

Answer: Seasonal Mann-Kendall is most appropriate because it is the only monotonic trend test that accounts for seasonality. Monthly averages can show how ozone changes by the months in each season.


13. Create a plot depicting mean monthly ozone concentrations over time, with both a `geom_point` and a `geom_line` layer. Edit your axis labels accordingly.

13

```
monthlyplot <- ggplot(GaringerOzone.monthly, aes(x=Date, y=mean_concentration))+
  geom_point() + geom_line () + xlab("Time") + ylab("Mean Ozone Concentration") +
  geom_smooth(method=lm)
```

```
print(monthlyplot)
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



14. To accompany your graph, summarize your results in context of the research question. Include output from the statistical test in parentheses at the end of your sentence. Feel free to use multiple sentences in your interpretation.

Answer: Ozone has been slightly decreasing over time at this station. This result is significant. Ozone is higher in the summer months than the winter months. ($\tau = -0.143$, 2-sided pvalue = 0.046724).

15. Subtract the seasonal component from the `GaringerOzone.monthly.ts`. Hint: Look at how we extracted the series components for the `EnoDischarge` on the lesson Rmd file.
16. Run the Mann Kendall test on the non-seasonal Ozone monthly series. Compare the results with the ones obtained with the Seasonal Mann Kendall on the complete series.

```
#15
MonthlyComponents <- as.data.frame(Monthlydecompose$time.series[,1:3])

MonthlyComponents <- mutate(MonthlyComponents,
  Mean_Ozone = GaringerOzone.monthly$mean_concentration,
```

```

Date = GaringerOzone.monthly$Date)

Garinger_unseas <- GaringerOzone.monthly.ts - Monthlydecompose$time.series[,1]
Monthly_Components <- mutate(MonthlyComponents,
Nonseasonal = Garinger_unseas)

#16

monthly_trend2 <- Kendall::MannKendall(Garinger_unseas)
summary(monthly_trend2)

## Score = -1179 , Var(Score) = 194365.7
## denominator = 7139.5
## tau = -0.165, 2-sided pvalue =0.0075402

```

Answer: These results had a smaller p-value than the seasonal Mann Kendall test. The tau is also slightly larger, indicating a slightly more negative trend. The results for both tests were statistically significant.